

The AI Evidence Checklist

Six things the second line is asked to produce when AI is challenged by a regulator, auditor, client or board. For each, the question is not whether you have a policy. It is whether you can produce the evidence on the day.

- ☐ 1 **What AI is in use, and who approved it**
PRODUCIBLE PROOF
A current AI inventory with named owner and approval record per tool. Not a wiki page from last quarter.
- ☐ 2 **Which customer or employee data can enter each tool**
PRODUCIBLE PROOF
A data-class map per tool: what is permitted in, what is blocked, enforced at runtime not by policy PDF.
- ☐ 3 **Which policy applies to each tool**
PRODUCIBLE PROOF
A policy-to-tool mapping a reviewer can follow without interpretation. Version-controlled.
- ☐ 4 **Which controls are logged, and where the log lives**
PRODUCIBLE PROOF
An access and usage log you can query: who used which model, when, on what. Retained 90 days minimum.
- ☐ 5 **Who was trained on the real failure mode**
PRODUCIBLE PROOF
A training record tied to the actual control, not a completion certificate for a generic module.
- ☐ 6 **What you would hand a regulator on the day they ask**
PRODUCIBLE PROOF
A reconstructable evidence chain for any single AI decision in under one hour. Prompt, model, cost, approval.

Score yourself

See your score in five minutes:
controller.intelxview.com/readiness